Improving Neural Network Performance

1. Overfitting
   - Dropout Layers
   - Regularization (L1 & L2)
   - Early Stopping

2. Normalization
   - Normalizing inputs
   - Batch Normalization
   - Normalizing Activations

3. Vanishing Gradients
   - Activation Functions
   - Weight Initializations

4. Gradient Checking & Clipping

5. Optimizers
   - Momentum
   - Adagard
   - RMSprop
   - Adam

6. Learning Rate Scheduling

7. Hyperparameter Tuning
   - No of hidden layers
   - Nodes in a layer
   - Batch Size
   - Epochs